

# 4M17 Coursework Assignment 2

## Practical Optimisation

Candidate Number: 5746G

January 10, 2021

### 1 Introduction

A general *optimisation problem* has the form:

$$\begin{aligned} & \text{minimize} && f_0(\mathbf{x}) \\ & \text{subject to} && f_i(\mathbf{x}) \leq b_i, \quad i = 1, \dots, m \end{aligned} \tag{1}$$

where the vector  $\mathbf{x} \in \mathbb{R}^n$  is the *n-dimensional optimisation variable* of the problem, the function  $f_0 : \mathbb{R}^n \rightarrow \mathbb{R}$  is the *objective function*, and the functions  $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ , for  $i = 1, \dots, m$ , are the *constraint functions*. The mathematical procedure of optimisation involves finding the *optimal solution*  $\mathbf{x}^*$  which has the smallest object value among all vectors which satisfy the constraints. That is to say, for any  $\mathbf{x}$  with  $f_i(\mathbf{x}) \leq b_i$  for  $i = 1, \dots, m$ , then  $f_0(\mathbf{x}) \geq f_0(\mathbf{x}^*)$  (Boyd & Vandenberghe 2004).

The first of the two 4M17 coursework assignments addressed the topic of *convex optimisation*, a subset of optimisation problems in which the objective and constraint functions are convex, meaning they satisfy the inequality:

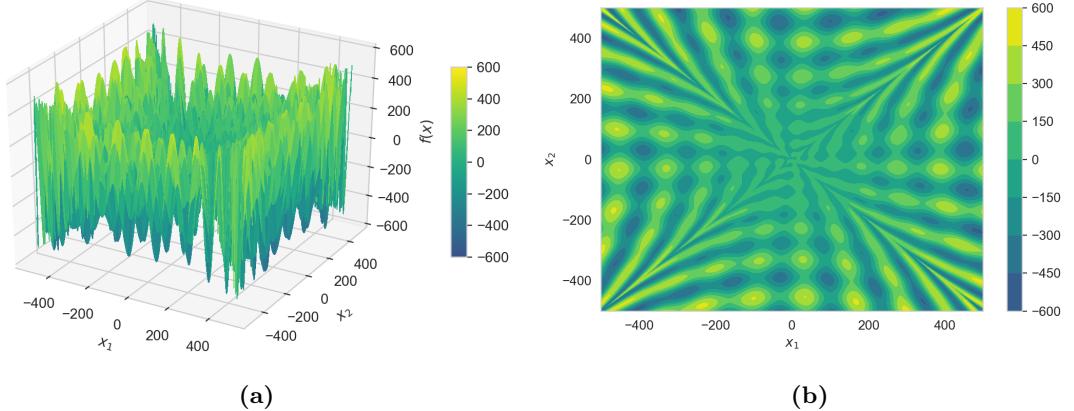
$$f_i(\alpha\mathbf{x} + \beta\mathbf{y}) \leq \alpha f_i(\mathbf{x}) + \beta f_i(\mathbf{y}), \tag{2}$$

for all  $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ , and all  $\alpha, \beta \in \mathbb{R}$  with  $\alpha + \beta = 1$ ,  $\alpha, \beta \geq 0$ . The advantage of working with convex problems is that they can be solved very efficiently and reliably using established numerical techniques, such as interior-point methods, often yielding polynomial solve times (Nesterov & Nemirovskii 1994). However, many real world optimisation problems are non-convex, and as such require a different category of optimisation techniques.

This report investigates two algorithms belonging to the class of *stochastic optimisation methods*: *simulated annealing* (SA) and *evolutionary strategies* (ESs). Both of these methods draw inspiration from real-world processes. There is a deep connection between SA and the process of *annealing* in metallurgy, a technique in which a metal is heated and its cooling controlled to produce a state of high order and very low energy (Kirkpatrick et al. 1982). ESs are a class of probabilistic search algorithms based on the *Darwinian model of evolution* (Darwin 1859, Schwefel 1977).

These heuristic approaches provide a means of iteratively searching over the search-space more efficiently than a purely *exhaustive search*. Probabilistic in nature, they introduce a degree of randomness into the search-process to accelerate the procedure of solving highly *multimodal nonlinear* mathematical optimisation problems. *Deterministic gradient methods* alone are unsuitable for problems with multimodalities, becoming 'trapped' in *local minima*. Stochastic algorithms, on the other hand, accept some changes which increase the objective function, without the need for first or second-order *derivative information*, thus overcoming this issue.

In general, these algorithms guarantee only to attain a locally optimum solution, which in some cases may be the *global optimum* (Schneider & Kirkpatrick, 2006). The process of *tuning* 'out of the box' algorithms to a problem specific function is explored, and the effects of varying parameter values on each algorithm discussed.



**Figure 1:** Panel (a) plots surface of the surface of Rana’s Function for  $n = 2$ . Panel (b) shows a visualisation of the contours, evenly spaced at intervals of 150.

## 2 The Problem

The problem investigated is *Rana’s Function* (Whitley et al. [1995, 1996]), an  $n$ -dimensional multi-modal optimisation problem with inequality constraints:

$$\begin{aligned} \text{minimize} \quad & f(\mathbf{x}) = \sum_{i=1}^{n-1} x_i \cos\left(\sqrt{|x_{i+1} + x_i + 1|}\right) \sin\left(\sqrt{|x_{i+1} - x_i + 1|}\right) \\ & + (1 + x_{i+1}) \cos\left(\sqrt{|x_{i+1} - x_i + 1|}\right) \sin\left(\sqrt{|x_{i+1} + x_i + 1|}\right) \\ \text{subject to} \quad & -500 \leq x_i \leq 500, \quad i = 1, \dots, n \end{aligned} \quad (3)$$

The report assesses the performance of the two stochastic algorithms on the 5-dimensional version of Rana’s Function (5D-RF), with one further constraint that only 10,000 objective function evaluations are permitted for each execution of an algorithm.

Figures 1(a) and 1(b) visualise the surface and contours of the lower dimensional form of the problem for  $n = 2$  (2D-RF). These plots reveal the extreme multimodality of the problem: the search-space is loitered with many local maxima and minima, resulting in the function being difficult in general to optimise.

Both plots in Figure 1 are generated using a *mesh-spacing* of 0.05, sub-dividing the input domain for each  $x_i \in [-500, 500]$  into  $500/0.05 = 10000$  discrete points. An *exhaustive search*, in which the 2D-RF is every evaluated point in the discretised search-space, requires  $10000^2$  evaluations of the 2D-RF. In general, for a mesh-spacing  $k$ , the domain of each  $x_i$  is sub-divided into  $500/k$  discrete points. Given an  $n$ -dimensional optimisation problem, this brute-force method requires the objective function to be evaluated at  $(500/k)^n$  points. The computational cost quickly becomes infeasible for increasing dimensionality; the number of iterations scales exponentially with the number of dimensions  $n$ . This undesirable trait, often referred to as ‘the curse of dimensionality’ (Taylor [1993]), motivates the need for efficient, scalable search algorithms.

The 2D-RF is the only form of Rana’s Function which can easily be visualised. It, therefore, provides a useful aid in representing search patterns followed by implementations of optimisation algorithms, in order to confirm they are operating as intended before being trialled on the higher dimensional 5D-RF.

### 3 Simulated Annealing

#### 3.1 Outline of Implementation

A version of the simulated annealing (SA) algorithm was implemented in Python, code for which is given in Appendix A.1.

An *initial solution*  $\mathbf{x}_0$  is generated by drawing from a uniform distribution over the entire search-space. Based on the suggestion by Kirkpatrick (1984), the initial temperature  $T_0$  is calculated from an initial search to give an average probability  $\chi_0 = 0.8$  of a solution that increases the objective function being accepted. Trial solutions are generated according to the formulation suggested by Parks (1990):

$$\mathbf{x}_{i+1} = \mathbf{x}_i + D\mathbf{u}, \quad (4)$$

where  $\mathbf{u}$  is a vector of uniform random variables over  $[-1, 1]$ , and  $D$  is a diagonal matrix which specifies the maximum change allowed in each variable. Proposed solutions  $\mathbf{x}_{i+1}$  which violate the inequality constraints of the optimisation problem in (3) are rejected immediately.

The elements of  $D$  are updated following each accepted changed according to the update rule:

$$D_{j+1} = (1 - \alpha)D_j + \alpha\omega R, \quad (5)$$

where the entries of the diagonal matrix  $R$  encode the magnitudes of the successful made to each element of  $\mathbf{x}_i$ :

$$R_{kk} = |D_{kk}u_k|, \quad (6)$$

and the *damping constant*  $\alpha$  determines the rate at which update information is incorporated into  $D$  with *weighting*  $\omega$ . Upper and lower limits  $D_u$  and  $D_l$  are placed on the values of the elements of  $D$  to prevent single values dominating or shrinking.

The *actual step size*  $\bar{d}$  is given by:

$$\bar{d}^2 = \sum_{k=1}^n R_{kk}^2, \quad (7)$$

from which the probability of accepting the trial solution is calculated as:

$$p = \max \left\{ 1, \exp \left( -\frac{\delta f}{T\bar{d}} \right) \right\}, \quad (8)$$

meaning all decreases in the objective function are accepted, and any increases accepted with probability  $\exp \left( -\frac{\delta f}{T\bar{d}} \right)$ , provided the trial solution  $\mathbf{x}_{i+1}$  is in the *feasible region*. Following the acceptance of a new solution, the *archive* of best solutions is updated. There is a maximum specified length  $L_k$  for the Markov chain at every temperature, from which a minimum number of acceptances is calculated as  $\eta_{min} = 0.6L$ . The temperature of the scheme is decremented following either  $L_k$  trials or  $\eta_{min}$  acceptances at a given temperature:

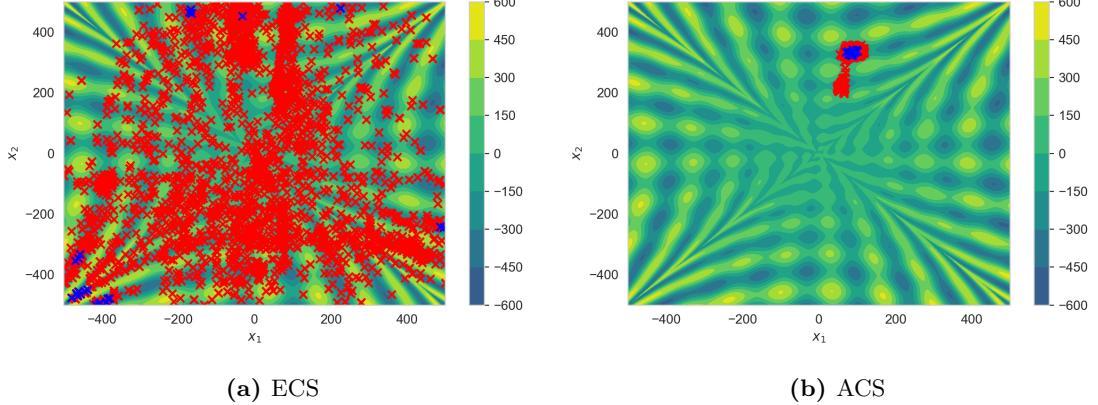
$$T_{k+1} = \alpha_k T_k, \quad (9)$$

where the factor  $\alpha_k = 0.95$  for the *exponential cooling scheme* (ECS) proposed by Kirkpatrick et al. (1982), or:

$$\alpha_k = \max \left\{ 0.5, \exp \left( -\frac{0.7T_k}{\sigma_k} \right) \right\} \quad (10)$$

for the *adaptive cooling scheme* (ACS) outlined by Huang et al. (1986), where  $\sigma_k$  is the standard deviation of the objective function values accepted at temperature  $T_k$ .

A *restart* is performed following a specified number of iterations over which no new best solution is found. The algorithm is *terminated* after reaching the maximum number of permitted of objective function evaluations or ceasing to make progress, whichever comes first. The latter condition is defined as no new best solution being found in an entire Markov chain for a given temperature, alongside the *solution acceptance ratio* falling below a specified threshold  $\chi_f = 0.01$ .



**Figure 2:** All accepted SA solutions (red) and final archived solutions (blue) for the ECS and ACS with  $L_k = 100$ .

### 3.2 Cooling Scheme and Maximum Markov Chain Length

The first parameters investigated were the choice of cooling<sup>1</sup> scheme, either exponential or adaptive, in conjunction with the maximum Markov chain length  $L_k$ . The temperature governs the balance between *exploration* of the search-space, and *exploitation* of good solutions. As the size of the temperature decrements is decreased, the probability that the algorithm terminates at the global optimum tends to 1 (Granville et al. 1994). However, infinitesimally small step sizes are practically infeasible, since the run-time of the algorithm also increases and quickly exceeds that of a purely exhaustive search. On the other hand, cooling too fast results in low probabilities of accepting steps which increase the objective function, increasing the likelihood of the optimiser becoming trapped in a local optimum. Thus, SA performance is conditional on selecting an appropriate annealing schedule. Finding an effective temperature decrement rule is often a problem specific task.

The reason for analysing the two parameters together is due to their interaction: the temperature is decremented after  $L_k$  iterations at the current temperature or after  $\eta_{min} = 0.6L_k$  acceptances, whichever occurs first. A sufficient number of iterations are required at each temperature such that a *quasi-equilibrium* state is achieved. This state is defined as the probability distribution of solutions  $q(N, T_k)$  after  $N$  trials at temperature  $T_k$  being ‘sufficiently close’ to the stationary distribution  $\pi(T_k)$  at  $T_k$ :

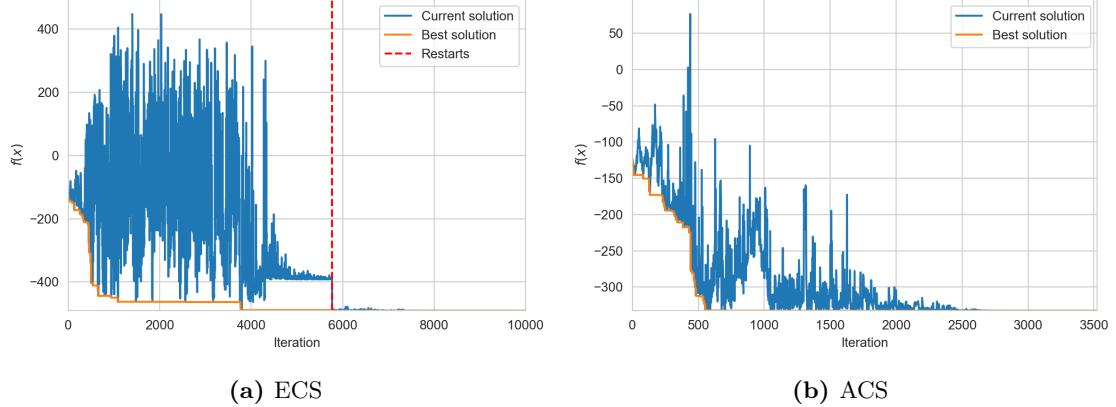
$$\|q(N, T_k) - \pi(T_k)\| < \epsilon, \quad (11)$$

where  $\epsilon$  is a small, positive value. It has been shown by Aarts et al. (1997) that larger drops in temperature require longer-length Markov chains to reach the quasi-equilibrium state, and so require larger values of  $L_k$ . Hence, the optimal configuration of  $L_k$  for a given problem is dependent on the choice of cooling scheme. This notion is highlighted by Figure 2, which plots all accepted solutions and the final archived solutions for both the ECS and ACS. Despite using the same maximum chain length of  $L_k = 100$ , a very different distribution of solutions is obtained for each of the cooling schemes.

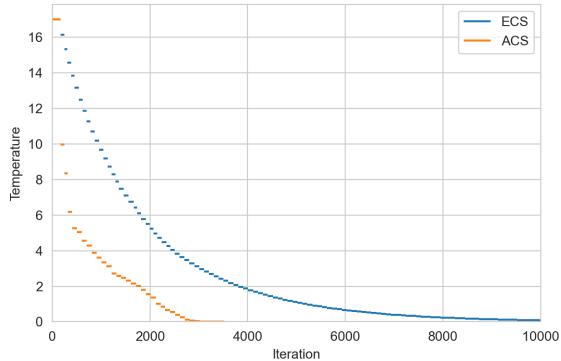
For the ECS, the annealer explored a significant portion of the search-space, solutions being generated for a wide range of  $(x_1, x_2)$  pairs. The dense bands of solutions, one between  $x_1 = 0$  and  $x_1 = 200$  spanning over the range of input values  $x_2 > 0$ , and another between  $x_2 = -200$  and  $x_2 = -400$  spanning over the range of input values  $x_1 > 0$ , correspond to samples accepted early on in the annealing process when the temperature was high. As the temperature cooled, the majority of samples were generated in the vicinity of the local minima, the algorithm settling on the regions where the function takes its lowest values. The 25 final archived solutions are positioned in similar such regions. This behaviour is verified by the variation of the evaluated objective function with SA

---

<sup>1</sup>In this report, the phrases *cooling* and *annealing* are used interchangeably. Both terms refer to the reduction of the temperature over the SA search.



**Figure 3:** Evolution of the objective function as a function of SA iterations for both candidate cooling schemes, ECS and ACS, corresponding to the solutions obtained in Figures 2(a) and 2(b), respectively.

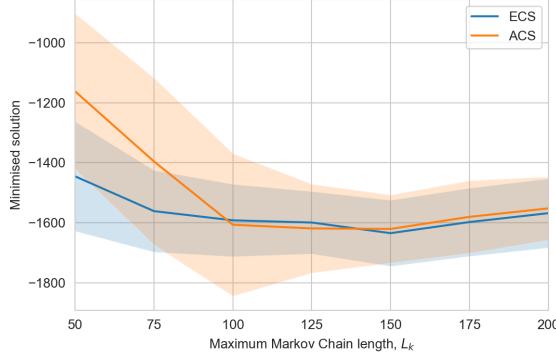


**Figure 4:** Variation of temperature with SA iterations for both candidate cooling schemes, ECS and ACS, corresponding to the solutions obtained in Figures 2(a) and 2(b), respectively.

iterations, a plot of which is shown in Figure 3(a). Early on in the process, when the temperature was high, the optimiser accepted a large number of solutions which increased the objective function, allowing it to navigate over the search-space. As the temperature was cooled, fewer objective function increases were accepted, and those which were took smaller magnitude. This is the typical search pattern taken by a functional SA algorithm, confirming the written code operated correctly on the 2D-RF.

In contrast to the ECS, the SA algorithm with the ACS explored only a very small proportion of the 2-dimensional search-space. Figure 4 plots the system temperature against SA iterations for both schemes, revealing that the ACS cooled significantly faster than the ECS. The search pattern of the adaptive annealing schedule is explained by the formulation for  $\alpha_k$  given in eq. (10): beginning in the vicinity of a local optimum, the variance  $\sigma_k^2$  of objective function values accepted was low for temperatures  $T_k$  early on in the optimisation procedure. This resulted in small corresponding values of  $\alpha_k$ , and hence rapid initial cooling. This meant increases in the objective function were assigned very small probability early on in the search routine. The annealer was therefore unable to escape the second local minimum it found, and eventually terminated based on the minimum solution acceptance criterion; the acceptance ratio for the 36<sup>th</sup> Markov chain dropped below  $\chi_f$ .

The effect of varying  $L_k$  for both cooling schemes was explored by performing 50 runs of the SA algorithm on the 5D-RF using different random number sequences. This was achieved by specifying 50 different *seeds* to the random number generator. Figure 5 plots the value of the best objective found averaged over the 50 runs. The uncertainty bars correspond to the mean plus and minus



**Figure 5:** Variation of the minimised 5D-RF for the exponential and adaptive cooling schedules with  $L_k$  averaged over 50 runs. The shaded area represents the mean plus and minus the standard deviation.

the standard deviation in the value of the best objective found. For the ECS, their width remains fairly constant, highlighting that the *consistency* of the annealing schedule is similar over the range of  $L_k$ 's tested. The width of the uncertainty bars for the ACS is large for lower values of  $L_k$ ; the high standard deviation indicates that some runs might yield potentially very good performance, but others quite poor performance. For both the exponential and adaptive cooling schemes, the lowest average best solution occurred for a maximum chain length of  $L_k = 150$ . With a lower overall variance and fractionally lower mean minimised solution, the exponential cooling scheme with  $L_k = 150$  was selected as the optimal annealing schedule for the minimisation of the 5D-RF.

### 3.3 Generation of Trial Solutions

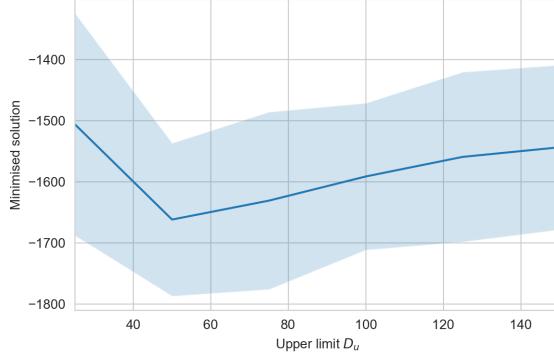
The next set of parameters investigated were those relating to the diagonal matrix  $D$ , which scales the elements of the vector of random variables  $\mathbf{u}$  in the generation of trial solutions, as given by eq. (4). These parameters are of particular importance, as they govern the distribution of obtainable solutions.

#### 3.3.1 Matrix Upper Limit

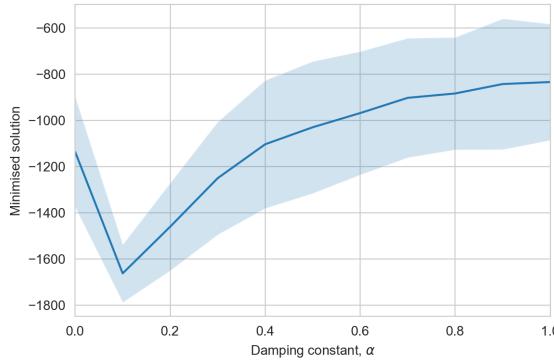
Figure 6 plots the value of the best objective found averaged over 50 runs as a function of the upper limit  $D_u$  imposed on the matrix  $D$ . The average best solution is minimised for a value  $D_u = 50$ . Lower values of  $D_u$  restrict the magnitude of the maximum permissible change in each variable of  $\mathbf{x}_i$ . The mobility of the annealer is inhibited, and it struggles to explore the entirety of the search-space. Therefore, it focuses itself on local minima encountered early on in the optimisation process, and often fails to reach the global minima. Conversely, higher values of  $D_u$  allow for larger step sizes. Individual values of  $D_u$  can grow larger, and the SA algorithm is able to traverse a greater portion of the search-space, even at low temperatures. As a result, its ability to refine the search to regions where the objective function takes lower values is hampered, resulting in larger average minimised objective function values.

#### 3.3.2 Damping Constant

From eq. (5), the damping constant  $\alpha$  determines the rate at which information from the matrix  $R$  is folded into the matrix  $D$ . The variation of the best objective found with  $\alpha$ , averaged over 50 runs, is plot in Figure 7. The average best objective is minimised for a value  $\alpha = 0.1$ . Higher values of  $\alpha$  incorporate information regarding the magnitude of the successful changes at a faster rate, resulting in the elements of  $D$  being increased more quickly. Since the elements of  $D$  specify



**Figure 6:** Minimised 5D-RF as a function of the upper limit  $D_u$ , averaged over 50 runs. The shaded area represents the mean plus and minus the standard deviation.



**Figure 7:** Variation of the value of the minimised 5D-RF as a function of the damping constant  $\alpha$ , averaged over 50 runs. The shaded area represents the mean plus and minus the standard deviation.

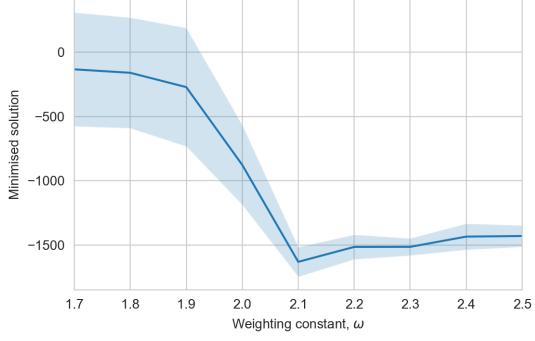
the maximum change allowed in each variable, larger elements of  $D$  allow for larger subsequent step sizes and hence greater exploration of the search-space. The converse is true for smaller  $\alpha$ , in which the elements of  $D$  are updated at a slower rate, leading to greater exploitation of good solutions. The optimal parameter setting  $\alpha = 0.1$  provides a trade-off between these two traits, a condition required for an effective search.

### 3.3.3 Weighting Constant

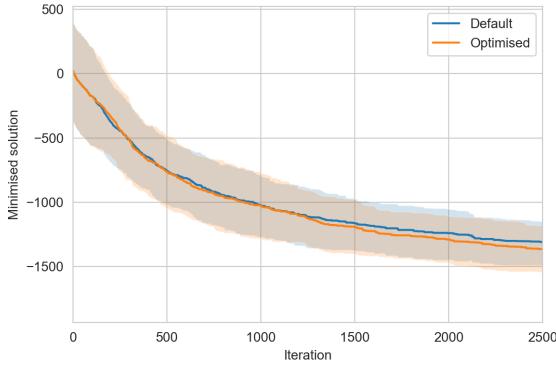
Closely tied to the damping constant is the weighting constant  $\omega$ , which from eq. (5) can be seen to control the weighting with which information from  $R$  is folded into  $D$ . Figure 8 plots the best objective found as  $\omega$  was varied, with  $\alpha$  fixed to its optimal value of 0.1, averaged over 50 runs. The average best objective was minimised for a value  $\omega = 2.1$ . This pair of constants  $(\alpha, \omega) = (0.1, 2.1)$  is consistent with the findings made by Parks (1990): this configuration tunes the maximum step size towards a value which yields acceptable changes as the search progresses.

## 3.4 Final Configuration

To evaluate the effect made by setting the investigated parameters to their optimal values, the minimised solution of the 5D-RF was recorded as a function of SA iterations for both the default configuration and the final, optimised configuration. Figure 9 plots the minimised solution averaged over 50 runs for both configurations. There was a significant performance gain obtained in the value of the minimised 5D-RF after 2500 iterations: on average, the optimal configuration returned



**Figure 8:** Variation of the value of the minimised 5D-RF as a function of the weighting constant  $\omega$ , averaged over 50 runs. The shaded area represents the mean plus and minus the standard deviation.



**Figure 9:** Evolution of the value of the minimised 5D-RF with SA iterations for both the default and optimal parameter configurations, averaged over 50 runs. The shaded area represents the mean plus and minus the standard deviation.

a minimised solution 5.88% lower than that of the default configuration. One potential reason why the difference was not greater is that many of the default solution generation parameters were set close to or equal to their optimal settings; only the cooling scheme, maximum Markov chain length  $L_k$  and matrix upper limit  $D_u$  parameters were modified. The uncertainty in the minimised solution for the optimal configuration, represented by the shaded area of the orange plot in Figure 9, was on average 5.40% larger than that of the default configuration, indicating that the optimal configuration performs less consistently than the default. In the majority of optimisation situations, the improvement in mean minimised solution for the 5D-RF is likely to outweigh the small increase in the standard deviation of the solutions, making the optimal configuration the better choice of parameter setting.

## 4 Evolution Strategies

### 4.1 Outline of Implementation

A Python implementation of an Evolution Strategy (ES), based primarily on the recommendations presented by Schwefel (1981), was devised, the code for which is presented in Appendix A.2. Where possible, explicit loops were avoided, and the code written using NumPy array expressions. By *vectorising* the code in this way, operations are applied simultaneously over many entries, in contrast to one entry at a time (as for a loop). Such parallelism yields significant cpu time performance gains:

vectorised array operations run one to two orders of magnitude faster than their pure Python loop counterparts (McKinney 2017).

Unlike Schwefel's original implementation, in which an *initial population* is generated by seeding an individual at a random point in the search-space and repeatedly applying *mutations*, an initial population is generated according to the method outlined by Bäck (1996): population members are conceived by sampling from an  $n$ -dimensional uniform distribution over the feasible region. The initial *covariance matrix* is generated by setting the diagonal variance elements  $\sigma_i^2(0) = 9.0$ , and the off-diagonal cross-covariance terms  $c_{ij}(0) = 0$ . Selecting small initial variances  $\sigma_i^2(0)$  prevents the algorithm diverging when the number of parents  $\mu$  is large (Bäck 1996). The operations in the algorithm work primarily in terms of *rotation angles*  $\alpha_{ij}$ , which are related to the covariances by:

$$\alpha_{ij} = \frac{1}{2} \tan^{-1} \left( \frac{2c_{ij}}{\sigma_i^2 - \sigma_j^2} \right). \quad (12)$$

In doing so, it is easier to guarantee the covariance matrix remains *positive semi-definite* (PSD).

For  $(\mu, \lambda)$ -selection, the  $\mu$  best individuals are selected *deterministically* from the pool of  $\lambda$  offspring. Whereas in  $(\mu + \lambda)$ -selection, the  $\mu$  best individuals are selected from the combined pool of  $\mu$  previous parents and their  $\lambda$  offspring. A new generation of  $\lambda$  offspring are created by first mutating the standard deviations and rotation angles (collectively referred to as the *strategy parameters*):

$$\sigma'_i = \sigma_i \exp(\tau' \times \aleph_0 + \tau \times \aleph_i) \quad (13)$$

$$\alpha'_{ij} = \alpha_{ij} + \beta \times \aleph_{ij} \quad (14)$$

where the  $\aleph$ s are realisations of normally distributed one dimensional random variables  $\mathcal{N}(0, 1)$  with mean zero and standard deviation one, and  $\tau$ ,  $\tau'$  and  $\beta$  are algorithm *control parameters* with values as suggested by Schwefel (1977):

$$\tau = \frac{1}{\sqrt{2\sqrt{n}}}; \quad \tau' = \frac{1}{\sqrt{2n}}; \quad \beta = \frac{5\pi}{180}.$$

The control variables are then mutated by adding a perturbation  $\mathbf{n}$ :

$$\mathbf{x}' = \mathbf{x} + \mathbf{n}, \quad (15)$$

where the vector  $\mathbf{n}$  is drawn from an  $n$ -dimensional Gaussian distribution  $\mathcal{N}(\mathbf{0}, C^{-1})$  with all-zero expectation vector and covariance matrix  $C^{-1} = \{c'_{ij}\}$ , the elements of which are given by:

$$c'_{ij} = \frac{1}{2} \left( \sigma'_i{}^2 - \sigma'_j{}^2 \right) \tan(2\alpha'_{ij}). \quad (16)$$

A small positive constant  $\epsilon_c$  is added to the diagonal elements of  $C^{-1}$ . Provided  $\epsilon_c$  is chosen large enough, this modified matrix is guaranteed to PSD, and hence a valid covariance matrix.

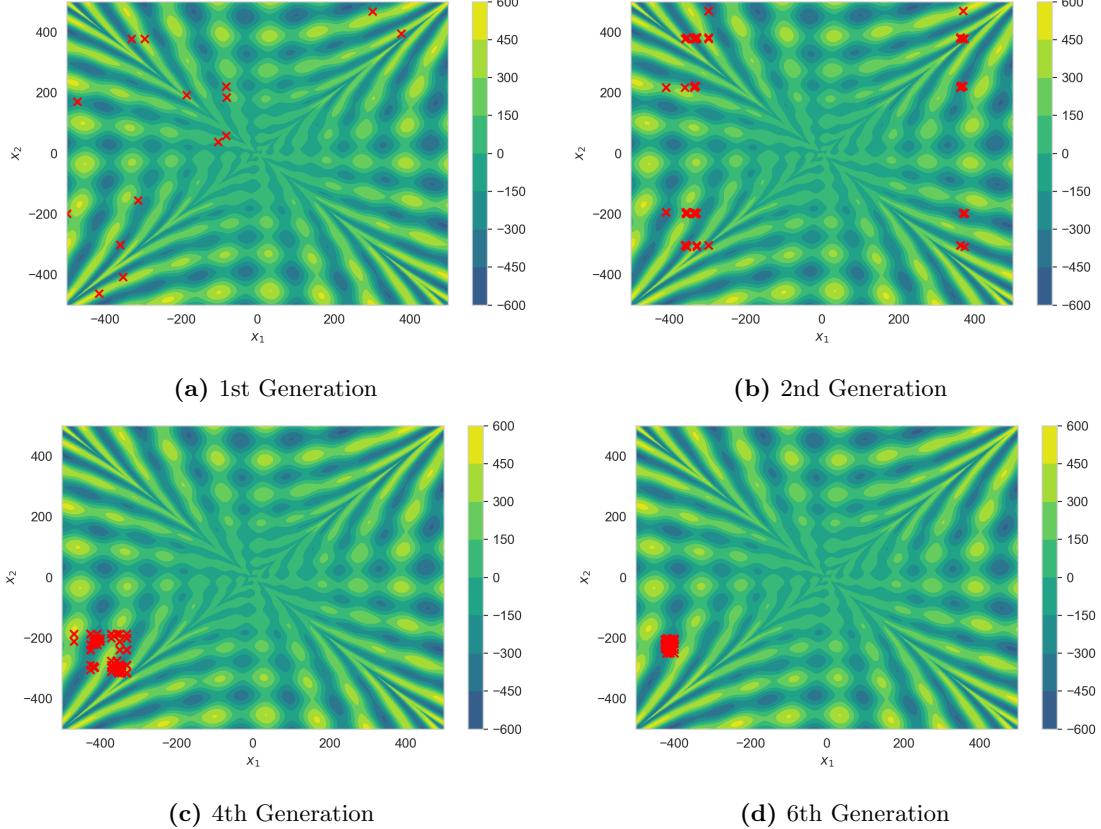
The multiplicative mutation process can result in individual standard deviations becoming too small. To prevent this, the algorithm constrains all standard deviation to a lower bound  $\epsilon_\sigma$ :

$$\sigma'_i < \epsilon_\sigma \Rightarrow \sigma_i := \epsilon_\sigma. \quad (17)$$

Rotation angles are kept in the feasible range of angle values  $[-\pi, \pi]$  by circularly mapping angles  $\alpha'_{ij}$  which exceed the limits:

$$|\alpha'_{ij}| > \pi \Rightarrow \alpha'_{ij} := \alpha'_{ij} - 2\pi \cdot \text{sign}(\alpha'_{ij}). \quad (18)$$

The  $\mu$  mutated control variables are *recombined* through *discrete recombination* to produce  $\lambda$  offspring solution. The component of each offspring solution is copied with probability 0.5 from each of its two randomly parent solutions. The strategy parameters are recombined using *intermediate*



**Figure 10:** Evolution of the population distribution for the minimisation of the 2D-RF using the ES algorithm with the default parameter configuration.

*recombination:* solution parameters are formed as a weighted average of the components from two randomly selected parents.

The inequality constraints of the problem are handled as in [Schwefel]’s original algorithm simply by removing offspring which violate the constraints, and repeating the creation of new solutions (through mutation and recombination) until the required number of offspring  $\lambda$  is obtained.

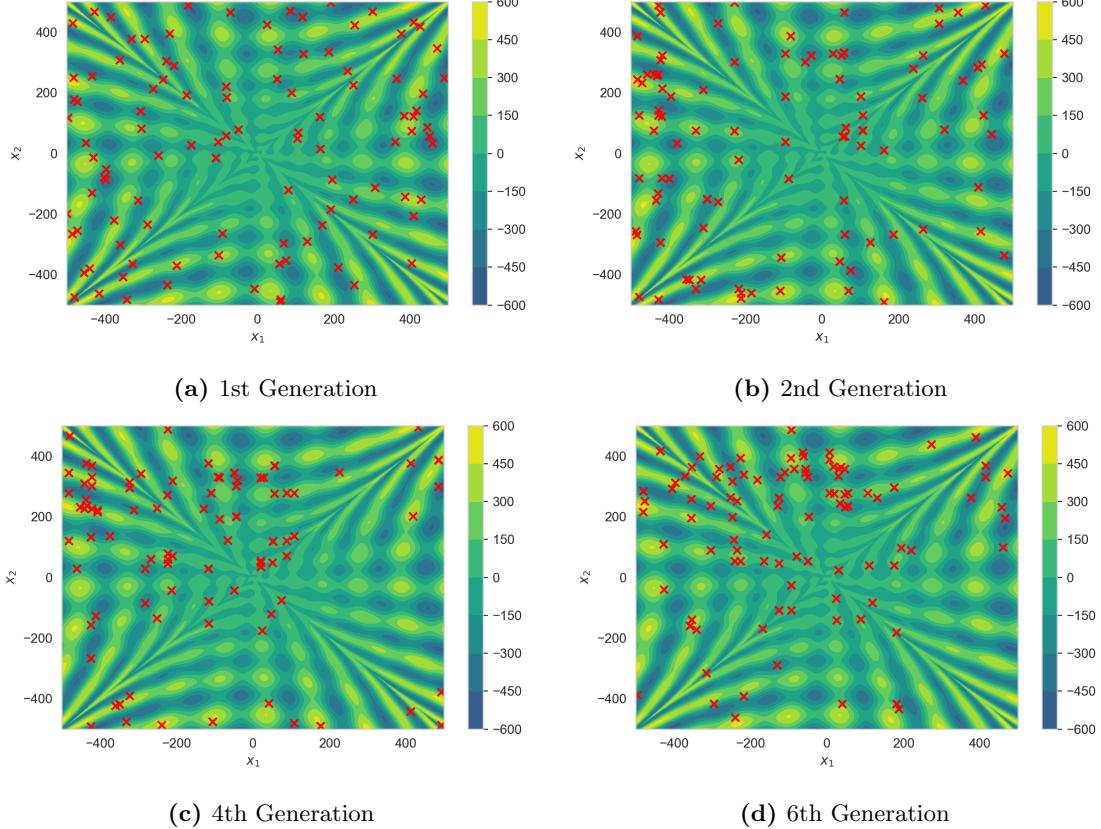
The search is terminated when the absolute difference between the worst and best solutions within a parent population falls below a user-specified threshold  $\epsilon > 0$ :

$$|f_w - f_b| < \epsilon \quad (19)$$

where  $f_w = \max_{\mathbf{x} \in \{\mathbf{x}_1 \dots \mathbf{x}_\mu\}} f(\mathbf{x})$  and  $f_b = \min_{\mathbf{x} \in \{\mathbf{x}_1 \dots \mathbf{x}_\mu\}} f(\mathbf{x})$ , or if the number of objective function evaluations exceeds the prescribed limit of 10,000.

## 4.2 2-Dimensional Rana Function

Figure 10 plots the population distribution obtained for a single run of the ES algorithm applied to the 2D-RF. The default parameter settings, as given in Appendix A.2, were used for this run. Each member of the 1st, 2nd, 4th and 6th generations is shown by a red cross. The variance in the spatial distribution of the population reduces as the algorithm progresses, converging to a local optimum at (-418, -220). A series of other runs are show in Appendix B. Together, they demonstrate the optimisation method worked as anticipated.



**Figure 11:** Evolution of the population distribution for the minimisation of the 2D-RF using the ES algorithm for  $(\mu, \lambda)$ -selection with the extreme case of  $\mu = \lambda = 15$ .

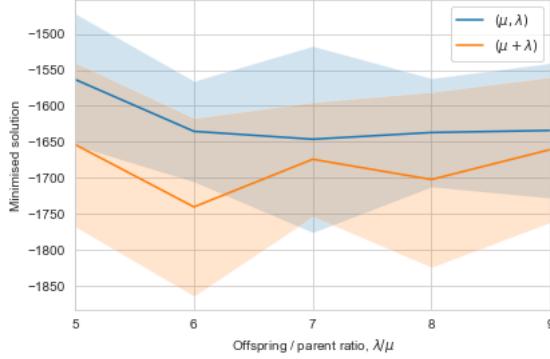
### 4.3 Selection

The selection operator provides the means by which ESs direct their search towards promising regions of the search space. At any given generation, selection exploits information represented within the population to reduce diversity by focusing on individual solutions where the evaluated objective function is lower.

There are two candidate selection schemes available for ESs: either  $(\mu, \lambda)$ , in which the  $\mu$  best individuals are selected from the  $\lambda$  offspring, or  $(\mu + \lambda)$ , in which the  $\mu$  best parents are selected from the union of the  $\mu$  previous parents and their  $\lambda$  offspring.

Each potential candidate schemes were trialled for different ratios of  $\mu : \lambda$ . To ensure  $(\mu, \lambda)$ -selection remained deterministic, the number of parents was kept lower than the number of offspring. In doing so, the random walk behaviour that occurs in the extreme case when  $\mu = \lambda$  was avoided. Figure 11 demonstrates this random walk behaviour for the 2-dimensional case; it plots snapshots of the population distribution for the minimisation of the 2D-RF using the same algorithm configurations as Figure 10, but with one change that  $\mu = \lambda$  instead of  $\mu = \lambda/7$ .

The performance of these potential schemes was investigated by applying the ES algorithm to the 5D-RF and recording the minimised solution. The number of parents was fixed as  $\mu = 15$ , and the ratio  $\lambda/\mu$  varied in integer intervals from  $\lambda/\mu = 5$  to  $\lambda/\mu = 9$ . The value of the best objective found for the 5D-RF averaged over 50 runs is plotted as a function of  $\lambda/\mu$  for both selection schemes in Figure 12. It shows that on average, the minimised solution obtained by the  $(\mu + \lambda)$  scheme is lower than that from the  $(\mu, \lambda)$ -selection over the range of ratios tested. Furthermore, the width of the error bars, corresponding to the mean plus and minus the standard deviation, remain tighter for  $(\mu + \lambda)$ -selection compared to  $(\mu, \lambda)$ . Thus for the 5D-RF, a  $(\mu + \lambda)$  selection scheme on average



**Figure 12:** Value of the minimised 5D-RF as a function of the offspring to parent ratio, averaged over 50 runs. The shaded area represents the mean plus and minus the standard deviation.

yields a lower minimised solution with greater consistency.

One potential reason for the superiority of  $(\mu + \lambda)$ -selection for the 5D-RF is the fact that it guarantees that the best individuals from the combined parent-offspring population survive to become the subsequent parents. Consequently, it produces a monotonic course of evolution, remembering good solutions from the previous reproduction cycle and carrying them forward to the next cycle. This is in contrast to  $(\mu, \lambda)$ -selection, in which individuals may only be selected in one reproduction cycle.

The best average solution found for the  $(\mu, \lambda)$  selection scheme occurred for a ratio of  $\mu : \lambda$  of 1 : 7. This is consistent with the findings of Schwefel [1987]. For  $(\mu + \lambda)$ -selection, a lower ratio of 1 : 6 was found to be optimal, returning on average the best overall solution for the two selection schemes trialled.

The ratio  $\mu : \lambda$  controls the degree of exploitation in the search, often termed the *selection pressure* (Bäck [1994]). Higher selection pressures place greater emphasis on better solutions, leading to increased exploitation. On the other hand, lower selection pressure enables worse solutions to survive, allowing for greater exploration. The optimal selection pressures provide a good balance between exploitation and exploration, a condition critical in order to achieve reasonable behaviour of an ES in the case of a complicated optimisation problem, such as the 5-dimensional Rana Function. Furthermore, selection pressures in this range are low enough to avoid *punctuated equilibrium* evolution (Gould & Eldredge [1977], Schwefel [1995]), an undesirable evolutionary phenomenon in which there are isolated episodes of rapid speciation amidst long periods of little or no change.

#### 4.4 Mutation

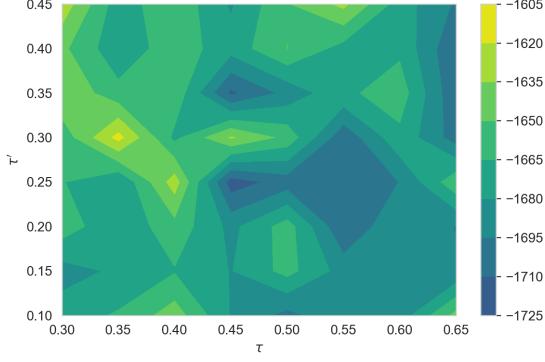
Mutation is the primary source of genetic variation in an ES algorithm. It serves the purpose of introducing *innovations*, a means by which the optimisation method explores the search-space (Bäck [1996], Schwefel [1995]).

From eq. (13), the control parameters  $\tau$  and  $\tau'$  dictate the standard deviations ('step sizes') of the  $\mathcal{N}(0, 1)$  normally distributed random variables  $N_i$  and  $N_0$ . These random variables are exponentiated, and then multiplied with  $\sigma_i$  to produce changes in the standard deviations. Similarly, from eq. (14), the parameter  $\beta$  governs the standard deviation of the perturbation applied to  $\alpha_{ij}$ .

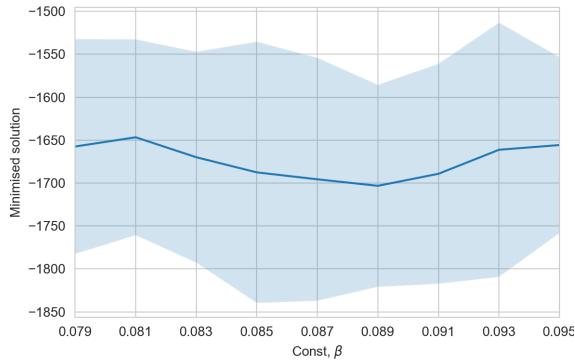
Under the recommendation of Schwefel [1977], by default the parameters take the following values when the ES algorithm is applied to the 5D-RF:

$$\tau = 0.473; \quad \tau' = 0.316; \quad \beta = 0.0873.$$

The effect of modifying the control parameters from their recommended values was investigated first by fixing  $\beta = 0.0873$  and jointly varying  $\tau$  and  $\tau'$ . The parameters  $\tau$  and  $\tau'$  were then fixed to their optimal values, and  $\beta$  varied.



**Figure 13:** Minimised value of the 5D-RF plot as a function of the control parameters  $\tau$  and  $\tau'$ , averaged over 50 runs. The parameter  $\beta$  was fixed to 0.0873.



**Figure 14:** Minimised value of the 5D-RF plot as a function of the control parameter  $\beta$ , averaged over 50 runs. The parameter pair  $(\tau, \tau')$  were fixed to their optimal values of  $(0.45, 0.25)$ .

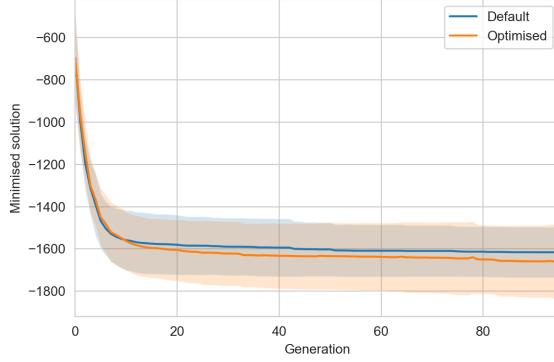
Figure 13 plots the contours of the minimised solution obtained when the ES was applied to the 5D-RF, averaged over 50 runs, as a function of  $\tau$  and  $\tau'$  with  $\beta$  held fixed. The lowest average solution occurred for a parameter pair  $(\tau, \tau') = (0.45, 0.25)$ . These two parameters were then fixed to their optimal values, and the minimised solution recorded as a function of  $\beta$ , a plot of which is shown in Figure 14. It shows that on average, the best minimised solution was obtained for  $\beta = 0.089$ , giving the trio of optimised control parameters:

$$\tau = 0.45; \quad \tau' = 0.25; \quad \beta = 0.089.$$

The first and last of these parameters take values close to those recommended by Schwefel (1977). However, the second parameter is significantly smaller, indicating a lower step size of individual-specific standard deviation mutations is more optimal for the ES applied to the 5D-RF.

## 4.5 Final Configuration

To quantify the effect of optimising the selection scheme and control parameters to the 5D-RF, the objective function was minimised using both the default and tuned ES algorithms, and the minimised solution recorded as a function of the ES generation. Figure 15 plots the evolution of the minimised solution with ES generation, averaged over 50 runs. On average, the final minimised solution for the optimised configuration was 5.43% lower than that of the default configuration, a significant performance gain. However, the uncertainty in the solution for the optimised settings was also marginally larger; the standard deviation for the tuned configuration was 3.32% greater its default counterpart. Together, these findings indicate that whilst on average the optimised parameter



**Figure 15:** Variation of the minimised solution for the ES algorithm applied to the 5D-RF averaged over 50 runs for both the default and optimised parameter configurations. The shaded area represents the mean plus and minus the standard deviation.

configuration yields better minimised objective values, it does so slightly less consistently. For most applications, the improvement in the mean minimised solution is likely to outweigh the small penalty in consistency, making the optimised configuration the better choice for algorithmic performance.

## 5 Conclusion

Whilst convex optimisation methods work efficiently and reliable when applied to convex problems, many real-world optimisation problems violate the constraints of strict convexity. Furthermore, deterministic gradient-based methods are inapplicable for problems involving multiple minima. For more complex problems, a different class of stochastic based optimisation methods are required.

The  $n$ -dimensional Rana-Function is one such example of a multimodal function. Constraining the feasible region to the domain  $x_i \in [-500, 500]$  for  $i = 1, \dots, n$ , this function was minimised using two different heuristic-based approaches, which were individually coded in Python. Critical to the performance of these methods is problem specific tuning of algorithmic parameters. A systematic investigation into a sub-set of parameters was performed for each of the methods presented.

The first algorithm explored was simulated annealing, for which there was shown to be a close tie-in between the cooling scheme and maximum Markov chain length. The exponential cooling scheme provided consistently lower minimised solutions compared to the adaptive scheme when applied to the 5D-RF. Modifying the parameters relating to the perturbation scaling matrix  $D$  effected the degree to which the annealer could traverse the search space. Tuning the damping and weighting parameters correctly was essential for providing maximum step sizes which yielded acceptable solution changes as the search progressed. The final optimised SA configuration produced a lower mean minimised solution than its default counterpart, with only a marginal decrease in algorithmic consistency.

The second search method focused on an evolution strategy implementation. A  $(\mu + \lambda)$  selection scheme produced better optimisation performance than  $(\mu, \lambda)$ -selection. The ratio of  $\mu : \lambda$  was set to provide a suitable level of selection pressure. The mutation control parameters were varied in a sequential manner, and optimal values found to give appropriate step sizes for mutation of the strategy parameters. On average, the final optimised configuration produced a lower minimised solution, with only a slight increase in the standard deviation of the best solution.

## References

- Aarts, E., Korst, J. & Laarhoven, van, P. (1997), *Simulated Annealing*, Wiley-Interscience.

- Bäck, T. (1994), Selective pressure in evolutionary algorithms: A characterization of selection mechanisms, in ‘Proceedings of the First IEEE Conference on Evolutionary Computation. IEEE World Congress on Computational Intelligence’, Vol. 1, pp. 57–62.
- Bäck, T. (1996), *Evolutionary Algorithms in Theory and Practice: Evolution Strategies, Evolutionary Programming, Genetic Algorithms*, Oxford University Press, Inc., USA.
- Boyd, S. & Vandenberghe, L. (2004), *Convex Optimization*, Cambridge University Press.
- Darwin, C. (1859), *On the Origin of Species by Means of Natural Selection*, Murray, London.
- Gould, S. J. & Eldredge, N. (1977), ‘Punctuated Equilibria: The Tempo and Mode of Evolution Reconsidered’, *Paleobiology* **3**(2), 115–151.
- Granville, V., Krivanek, M. & Rasson, J. . (1994), ‘Simulated annealing: a proof of convergence’, *IEEE Transactions on Pattern Analysis and Machine Intelligence* **16**, 652–656.
- Huang, M. D., Romeo, F. & Sangiovanni-Vincentelli (1986), ‘An Efficient General Cooling Schedule for Simulated Annealing’, *Proceedings of the IEEE International conference on Computer Aided Design* pp. 381–384.
- Kirkpatrick, S. (1984), ‘Optimization by Simulated Annealing: Quantitative Studies’, *Journal of Statistical Physics* **34**, 975–986.
- Kirkpatrick, S., Gelatt, C. D. & Vecchi, M. P. (1982), ‘Optimization by Simulated Annealing’, *IBM Research Report RC 9355*.
- McKinney, W. (2017), *Python for Data Analysis: Data Wrangling with Pandas, NumPy, and IPython*, 2nd edn, O’Reilly Media, Inc.
- Nesterov, Y. & Nemirovskii, A. (1994), *Interior Point Polynomial Algorithms in Convex Programming*, Society for Industrial and Applied Mathematics, Philadelphia, Pa.
- Parks, G. T. (1990), ‘An Intelligent Stochastic Optimization Routine for Nuclear Fuel Cycle Design’, *Nuclear Technology* **89**, 233–246.
- Schneider, J. J. & Kirkpatrick., S. (2006), *Stochastic optimization*, Scientific computation, Springer, Berlin.
- Schwefel, H. (1987), Collective Phenomena in Evolutionary Systems, in ‘Preprints of the 31st Annual Meeting’, International Society for General System Research, Budapest, pp. 1025–1033.
- Schwefel, H.-P. (1977), *Numerische Optimierung von Computermodellen mittels der Evolutionstrategie*, Vol. 26, Birkhäuser, Basel Stuttgart.
- Schwefel, H.-P. (1981), *Numerical Optimization of Computer Models*, John Wiley & Sons, Inc., USA.
- Schwefel, H.-P. (1995), *Evolution and Optimum Seeking*, John Wiley & Sons, Inc., USA.
- Taylor, C. R. (1993), *Applications of dynamic programming to agricultural decision problems*, Westview Press, Boulder.
- Whitley, D., Rana, S., Dzubera, J. & Mathias, K. E. (1996), ‘Evaluating Evolutionary Algorithms’, *Artificial Intelligence* **85**, 245 – 276.
- Whitley, L. D., Mathias, K. E., Rana, S. B. & Dzubera, J. (1995), Building Better Test Functions, in ‘Proceedings of the 6th International Conference on Genetic Algorithms’, Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, p. 239–247.

# Appendix

## A Code Listings

### A.1 Simulated Annealing

---

```
1 import numpy as np
2 from numpy.linalg import norm
3
4
5 def f(x):
6     """
7         Function which evaluates the Rana-Function for an n-dimensional solution.
8
9     Parameters
10    -----
11     x : (n, ) array
12         Solution to be evaluated.
13
14    Returns
15    -----
16    f_eval : float
17        Objective function evaluated at the input x.
18    """
19    # Initialise the output value
20    f_eval = 0
21    # Summation over (n-1) dimensions
22    for i in range(x.shape[0] - 1):
23        f_eval += x[i] * np.cos(np.sqrt(np.abs(x[i + 1] + x[i] + 1))) * np.sin(
24            np.sqrt(np.abs(x[i + 1] - x[i] + 1))) + (1 + x[i + 1]) * np.cos(
25            np.sqrt(np.abs(x[i + 1] - x[i] + 1))) * np.sin(np.sqrt(np.abs(x[i + 1]
+ x[i] + 1)))
26    return f_eval
27
28
29 class SA:
30     """
31         Simulated annealing optimisation to find the minimum of an objective function
32         f. The initial annealing temperature is found by performing an initial
33         search and selecting a temperature which yields a specified average
34         probability of acceptance for increases in f. The trial solution formula
35         suggested by Parks [1990] takes the form:
36          $x_{i+1} = x_i + Du$ ,
37     
```

where  $u \sim U(-1, 1)$ , and  $D$  is a diagonal matrix which specifies the maximum change permitted in each variable.  $D$  is updated after successful changes to the objective function by folding in information regarding the magnitudes of successful changes made. The temperature is decremented adaptively by the formulation from Huang et al. [1986] once a given number of trials have been performed at the current temperature, or if a minimum number of acceptances have been made, whichever comes first. If a given number of iterations have elapsed without improvement to the current best solution, a restart is performed, and the search resumed from the current best solution. The archive stores the current best solutions, based on a Euclidean distance dissimilarity criterion.

```

34
35
36     def __init__(self, f, n, bound, max_iter=10000, seed=0, max_chain=100,
37                  chi_0=0.8, alpha=0.1, omega=2.1,
38                  d_u=None, d_l=None, cooling_scheme='adaptive', max_restart=2000,
39                  chi_f=0.01, archive_size=25,
40                  d_min=None, d_sim=None):
41
42     """
43         Class attribute initialisation.
44
45     Parameters
46     -----
47
48     f : func
49         Objective function to be minimised.
50
51     n : int
52         Number of dimensions over which to optimise.
53
54     bound : float
55         Symmetric bound (+/-) on the domain of the control variables.
56
57     max_iter : int, optional
58         Maximum number of objective function evaluations. The default is
59         10000.
60
61     seed : int, optional
62         Random number generator seed for reproducibility. The default is 0.
63
64     max_chain : int, optional
65         Number of iterations after which to terminate Markov chain and
66         decrement temperature. The default is 100.
67
68     chi_0 : float, optional
69         Avg probability of accepting increases in f during initial search. The
70         default is 0.8.
71
72     alpha : float, optional
73         Damping constant. The default is 0.1.
74
75     omega : float, optional
76         Weighting constant. The default is 2.1.
77
78     d_u : float, optional
79         Upper limit on the elements of the diagonal matrix D. If d_u is None
80         (default), d_u is set to bound / 5.
81
82     d_l : float, optional
83         Lower limit on the elements of the diagonal matrix D. If d_l is None
84         (default), d_l is set to d_u / 1e5.
85
86     cooling_scheme : str, optional
87         Choice of cooling scheme, either 'exponential' (Kirkpatrick et al.
88         [1982]) or 'adaptive' (Huang et al. [1986]). The default is 'adaptive'.
89
90     max_restart : int, optional

```

```

69      Number of iterations after which to restart search from the current
70      best solution. The default is 2000.
71      chi_f : float, optional
72          Solution acceptance ratio termination threshold. The default is 0.01.
73      archive_size : int, optional
74          Number of solutions to store in the archive. The default is 25.
75      d_min : float, optional
76          Archive dissimilarity threshold. If d_min is None (default), d_min is
77          set to bound / 100.
78      d_sim : float, optional
79          Archive dissimilarity threshold. If d_sim is None (default), d_sim is
80          set to d_min / 100.
81      """
82      # assertion check
83      assert cooling_scheme in ['exponential', 'adaptive'], \
84          "cooling scheme must either be 'exponential' or 'adaptive'"
85      # initialise attributes
86      self.f = f
87      self.n = n
88      self.bound = np.abs(bound)
89      self.max_iter = max_iter
90      self.seed = seed
91      self.max_chain = max_chain
92      self.chi_0 = chi_0
93      self.alpha = alpha
94      self.omega = omega
95      if d_u:
96          self.d_u = d_u
97      else:
98          self.d_u = self.bound / 5
99      if d_l:
100         self.d_l = d_l
101     else:
102         self.d_l = self.d_u / 1e5
103     self.cooling_scheme = cooling_scheme
104     self.max_restart = max_restart
105     self.chi_f = chi_f
106     self.archive_size = archive_size
107     if d_min:
108         self.d_min = d_min
109     else:
110         self.d_min = self.bound / 100
111     if d_sim:
112         self.d_sim = d_sim
113     else:
114         self.d_sim = self.d_min / 100
115     # random seed for reproducibility
116     np.random.seed(seed=seed)
117     # generate initial solution by drawing from a uniform random variable over
118     # the feasible region
119     self.x0 = np.random.uniform(-bound, bound, n)

```

```

120     # generate non-zero elements of the diagonal matrix D, represented as a
121     # (n, ) vector d
122     self.d = np.ones(n)
123     # initialise iteration and acceptance counters
124     self.iter = 1
125     self.acc = 1
126
127     def initialise_T(self):
128         """
129             Method which initialises the search temperature  $T_0$  by performing an
130             initial search in which all increases in  $f$  are accepted.  $T_0$  is calculated
131             to give an average probability  $\chi_0$  of a solution that increases  $f$ .
132
133             Returns
134             -----
135             x_prev : (n, ) array
136                 Vector of current solution.
137             f_prev : float
138                 Objective function evaluated at current solution.
139             T_0 : float
140                 Initial search temperature.
141
142             # initial solution x0
143             x_prev = self.x0
144             # evaluate objective function at initial solution
145             f_prev = self.f_x0
146             # array of objective increases
147             df_plus = []
148
149             # run the initial search for specified maximum length of Markov chain
150             while self.iter < self.max_chain:
151                 # generate a uniform rv in the range (-1,1) of equal dimension as x
152                 u = np.random.uniform(-1, 1, self.n)
153                 # find the proposal step dx
154                 dx = self.d * u
155                 # perturb the current position x_prev by the proposal step dx
156                 x_new = x_prev + dx
157                 # reject proposal if outside of feasible region
158                 if max(abs(x_new)) > self.bound:
159                     # increment iteration counter
160                     self.iter += 1
161                     continue
162
163                 # evaluate change in objective function
164                 f_new = self.f(x_new)
165                 df = f_new - f_prev
166                 # record increases in objective function
167                 if df > 0:
168                     df_plus.append(df)
169                 # update d only if objective function is decreased(?)
170                 r = np.abs(dx)
171                 self.d = (1 - self.alpha) * self.d + self.alpha * self.omega * r
172                 self.d = np.clip(self.d, self.d_l, self.d_u)
173                 # accept all feasible trial solutions
174                 x_prev = x_new
175                 f_prev = f_new

```

```

172     # update archive
173     self.update_archive(x_prev, f_prev)
174     # increment iteration and acceptance counters
175     self.iter += 1
176     self.acc += 1
177
178     # average objective function increase
179     df_bar = np.mean(df_plus)
180     T_0 = - df_bar / np.log(self.chi_0)
181     return x_prev, f_prev, T_0
182
183 def update_T(self, T, f_T):
184     """
185         Method which updates the annealing temperature according to either: the
186         exponential CS outlined by Kirkpatrick et al. [1982]; the adaptive CS
187         outlined by Huang et al. [1986].
188
189     Parameters
190     -----
191     T : float
192         Current temperature.
193     f_T : (k, ) array
194         Vector of objective function acceptances at current temperature.
195
196     Returns
197     -----
198     alpha * T : float
199         Decremented temperature.
200     """
201
202     if self.cooling_scheme == 'exponential':
203         alpha = 0.95
204     else:
205         if len(f_T) > 1:
206             sigma = np.std(f_T)
207             alpha = max(0.5, np.exp(-0.7 * T / sigma))
208         else:
209             alpha = 0.95
210     return alpha * T
211
212 def update_archive(self, x_j, f_j):
213     """
214         Method which updates the class archives of best solutions.
215
216     Parameters
217     -----
218     x_j : (n, ) array
219         A new n-dimensional candidate solution for archiving.
220     f_j : float
221         Objective function evaluated at the candidate solution x_j.
222
223     # worst archived solution
224     g_ind = np.argmax(self.archive_f)
225     f_g = self.archive_f[g_ind]
226
227     # Euclidean distance between new and archived solutions

```

```

225     distance = norm(self.archive_x - x_j, axis=1)
226     # archived solution  $x_e$  which most closely resembles new solution  $x_j$ 
227     e_ind = np.argmin(distance)
228     d_ej = distance[e_ind]
229     f_e = self.archive_f[e_ind]
230
231     # archive is not full
232     if len(self.archive_x) < self.archive_size:
233         # archive new solution  $x_j$  if it is sufficiently dissimilar to all the
234         # solutions archived
235         if d_ej > self.d_min:
236             self.archive_x.append(x_j)
237             self.archive_f.append(f_j)
238
239     # archive is full
240     else:
241         # archive  $x_j$  if it is sufficiently dissimilar to all the solutions
242         # archived and better than the worst
243         if d_ej > self.d_min and f_j < f_g:
244             # new solution  $x_j$  replaces the worst archived solution  $x_g$ 
245             self.archive_x.pop(g_ind)
246             self.archive_f.pop(g_ind)
247             self.archive_x.append(x_j)
248             self.archive_f.append(f_j)
249             # archive  $x_j$  if it is not sufficiently dissimilar to all the
250             # solutions archived and better than the worst
251             if d_ej < self.d_min and f_j < f_g:
252                 # new solution  $x_j$  replaces the closest archived solution  $x_e$ 
253                 self.archive_x.pop(e_ind)
254                 self.archive_f.pop(e_ind)
255                 self.archive_x.append(x_j)
256                 self.archive_f.append(f_j)
257             # archive  $x_j$  if it is not the best solution so far and sufficiently
258             # similar to but better than  $x_e$ 
259             if d_ej < self.d_sim and f_j < f_e:
260                 self.archive_x.pop(e_ind)
261                 self.archive_f.pop(e_ind)
262                 self.archive_x.append(x_j)
263                 self.archive_f.append(f_j)
264
265     def run(self):
266         """
267             Method which runs the SA algorithm.
268
269             Returns
270             -----
271             x_min : (n, ) array
272                 Minimised solution.
273             f_min : float
274                 Objective function evaluated at the minimised solution.
275             archive_x : (k, n) array
276                 Matrix of k final archived solutions. Each control variable is of
277                 dimension n.
278             archive_fx : (k, ) array

```

```

274     Vector of objective function values evaluated at the k final archived
275     solutions.
276     """
277
278     # perform initial search
279     x_prev, f_prev, T = self.initialise_T()
280     # set minimum number of acceptances for each temperature
281     eta = 0.6 * self.max_chain
282     # initialise iterations and acceptances for current temperature
283     iter_T, acc_T = 0, 0
284     # initialise restart counter
285     restart_iter = 0
286     # array of objective function acceptances at current temperature
287     f_T = [f_prev]
288
289     # run the search routine for a given maximum number of iterations
290     while self.iter < self.max_iter:
291         # generate a uniform rv in the range (-1,1) of equal dimension as x
292         u = np.random.uniform(-1, 1, self.n)
293         # find the proposal step dx
294         dx = self.d * u
295         # generate the elements of the diagonal matrix R, represented as a
296         # vector r
297         r = np.abs(dx)
298         # compute the actual step size associated with the proposal step dx
299         step_size = np.sqrt(np.sum(r ** 2))
300         # perturb the current position x_prev by the proposal step dx
301         x_new = x_prev + dx
302         # increment iteration counters
303         self.iter += 1
304         iter_T += 1
305         restart_iter += 1
306
307         # accept proposal if inside feasible region
308         if max(abs(x_new)) < self.bound:
309             # evaluate objective function at proposal solution
310             f_new = self.f(x_new)
311             # acceptance probability
312             p_acc = min(1, np.exp(-(f_new - f_prev) / (T * step_size)))
313             if p_acc > np.random.uniform(0, 1):
314                 # reset restart counter if new best solution has been found
315                 if f_new < min(self.archive_f):
316                     restart_iter = 0
317                     # update archive
318                     self.update_archive(x_new, f_new)
319                     # update non-zero elements of matrix D
320                     self.d = (1 - self.alpha) * self.d + self.alpha * self.omega *
321                         r
322                     self.d = np.clip(self.d, self.d_l, self.d_u)
323                     # update array of objective function acceptances at current
324                     # temperature
325                     f_T.append(f_new)
326                     # accept solution
327                     x_prev = x_new
328                     f_prev = f_new
329                     # increment acceptance counters

```

```

325         acc_T += 1
326         self.acc += 1
327
328     # check conditions for termination
329     # terminate if no new best solution over whole Markov chain and
330     # acceptance ratio falls below threshold chi_f
331     if iter_T > self.max_chain and restart_iter > self.max_chain and acc_T
332         / iter_T < self.chi_f:
333         break
334
335     # check conditions for temperature cooling
336     # decrement temperature if max Markov chain length is attained or
337     # number acceptances exceeds eta
338     if iter_T > self.max_chain or acc_T > eta:
339         # standard deviation of objective function values at current
340         # temperature
341         T = self.update_T(T, f_T)
342         # reset temperature specific counters and array
343         iter_T, acc_T = 0, 0
344         f_T = []
345
346     # check conditions for restart
347     if restart_iter > self.max_restart:
348         # archive index of best solution
349         best_ind = np.argmin(self.archive_f)
350         # set current solution to best solution
351         x_prev = self.archive_x[best_ind]
352         f_prev = self.archive_f[best_ind]
353         # reset restart counter
354         restart_iter = 0
355
356     # best minimised solution
357     idx = np.argmin(self.archive_f)
358     x_min = self.archive_x[idx]
359     f_min = self.archive_f[idx]
360
361     return x_min, f_min, self.archive_x, self.archive_f
362
363
364
365
366
367 def main():
368     # number of dimensions and bound on feasible region
369     n, bound = 5, 500
370     sa = SA(f, n, bound)
371     x_min, f_min, archive_x, archive_f = sa.run()
372
373
374
375
376
377 if __name__ == '__main__':
378     main()

```

---

## A.2 Evolution Strategies

---

```

1 import numpy as np
2 from numpy.linalg import norm

```

```

3
4
5 def f(x):
6     """
7         Function which evaluates the Rana-Function for an array of n-dimensional
8             solutions.
9
10    Parameters
11    -----
12    x : (k, n) array
13        Input array of k control variables, each of dimension n.
14
15    Returns
16    -----
17    f_eval : (k, ) array
18        Output array of k objective function values.
19
20    # Initialise the array of output values
21    f_eval = np.zeros(x.shape[0])
22    # Summation over (n-1) dimensions
23    for i in range(x.shape[1] - 1):
24        f_eval += x[:, i] * np.cos(np.sqrt(np.abs(x[:, i + 1] + x[:, i] + 1))) *
25            np.sin(
26                np.sqrt(np.abs(x[:, i + 1] - x[:, i] + 1))) + (1 + x[:, i + 1]) *
27                    np.cos(
28                        np.sqrt(np.abs(x[:, i + 1] - x[:, i] + 1))) *
29                            np.sin(np.sqrt(np.abs(x[:, i + 1] + x[:, i] + 1)))
30
31    return f_eval
32
33
34 class ES:
35     """
36         Class for Evolutionary Strategy (ES) optimisation to find the minimum of an
37             objective function f. An initial population x is generated by sampling
38                 from a uniform distribution over the feasible region. Parents are
39                     deterministically selected, and then mutated through a stochastic process.
40                     The mutated parents are recombined to form offspring. The objective
41                         function f is then evaluated for the newly formed offspring. The
42                             evolutionary cycle of selection, mutation and recombination repeats. The
43                                 algorithm is terminated when the absolute difference in the objective
44                                     function of the best and worst members of a parent population falls within
45                                         a user-prescribed limit.
46
47     """
48
49     def __init__(self, f, n, bound, max_iter=10000, seed=0, mu=15, lambda_=None,
50                  sel=',', epsilon=1, epsilon_s=0.05,
51                  epsilon_c=1e-6, archive_size=25, d_min=None, d_sim=None,
52                  tau=None, tau_d=None, beta=None):
53
54         """
55             Class attribute initialisation.
56
57         Attributes
58         -----
59         f : func
60             Objective function to be minimised.

```

```

43     n : int
44         Number of dimensions over which to optimise.
45     bound : float
46         Symmetric bound (+/-) on the domain of the control variables.
47     max_iter : int, optional
48         Maximum number of objective function evaluations. The default is
49             10000.
50     seed : int, optional
51         Random number generator seed for reproducibility. The default is 0.
52     mu : int, optional
53         Number of parents. The default is 15.
54     lambd : int, optional
55         Number of offspring. If lambd is None (default), the number of
56         offspring is set to 7 * mu.
57     sel : str, optional
58         Selection scheme, either "," or "+". The default is ','
59     epsilon : float, optional
60         Termination threshold. The default is 1.
61     epsilon_s : float, optional
62         Lower limit on the standard deviation of the control variables. The
63         default is 0.05
64     epsilon_c : float, optional
65         Positive constant by which to scale the identity matrix I added to the
66         covariance matrix. Ensures the resulting matrix is PSD. The default is
67         1e-6.
68     tau : float, optional
69         Mutation control parameter. If tau is None (default), tau is set to 1
70         / sqrt(2 * sqrt(n)) (Schwefel [1977]).  

71     tau_d : float, optional
72         Mutation control parameter. If tau is None (default), tau_d is set to
73         1 / sqrt(2 * n) (Schwefel [1977]).  

74     beta: float, optional
75         Mutation control parameter. If beta is None (default), beta is set to
76         5 * pi / 180 (Schwefel [1977]).  

77     archive_size : int, optional
78         Number of solutions to store in the archive. The default is 25.
79     d_min : float, optional
80         Archive dissimilarity threshold. If d_min is None (default), d_min is
81         set to bound / 100.  

82     d_sim : float, optional
83         Archive dissimilarity threshold. If d_sim is None (default), d_sim is
84         set to d_min / 100.
85     """
86     # assertion check- selection scheme must be (mu, lambd) or (mu + lambd)
87     assert sel in [',', '+'], "selection scheme must either be ',' or '+'"
88     # initialise attributes
89     self.f = f
90     self.n = n
91     self.bound = np.abs(bound)
92     self.max_iter = max_iter
93     self.seed = seed
94     self.mu = mu
95     if lambd:
96         self.lambd = lambd
97     else:

```

```

88         self.lambd = 7 * self.mu
89         self.sel = sel
90         self.epsilon = epsilon
91         self.epsilon_s = epsilon_s
92         self.epsilon_c = epsilon_c
93         self.archive_size = archive_size
94         if d_min:
95             self.d_min = d_min
96         else:
97             self.d_min = self.bound / 100
98         if d_sim:
99             self.d_sim = d_sim
100        else:
101            self.d_sim = self.d_min / 100
102        if tau:
103            self.tau = tau
104        else:
105            self.tau = 1 / np.sqrt(2 * np.sqrt(self.n))
106        if tau_d:
107            self.tau_d = tau_d
108        else:
109            self.tau_d = 1 / np.sqrt(2 * self.n)
110        if beta:
111            self.beta = beta
112        else:
113            self.beta = 5 * np.pi / 180
114
115    def initialise(self):
116        """
117            Method which initialises the ES.
118
119            Returns
120            -----
121            xp : (mu, n)
122                Matrix of mu initial n-dimensional parent solutions.
123            fp : (mu, )
124                Vector of objective function values evaluated at the mu parent
125            solutions.
126            sigmap : (mu, n)
127                Matrix of parent standard deviations.
128            alphap : (mu, n)
129                Matrix of parent rotation angles.
130            iters : int
131                Number of objective function evaluations.
132        """
133
134        # random seed for reproducibility
135        np.random.seed(seed=self.seed)
136        # generate the initial population by sampling from a multivariate uniform
137        # distribution over the feasible region
138        xp = np.random.uniform(-self.bound, self.bound, [self.mu, self.n])
139        # evaluate the objective function
140        fxp = self.f(xp)

```

```

141     sigmap = 3 * np.ones((self.mu, self.n))
142     # initialise all rotation angles to 0
143     alphap = np.zeros((self.mu, self.n, self.n))
144     return xp, fxp, sigmap, alphap, iters
145
146 def selection(self, xp, fxp, sigmap, alphap, xo, fxo, sigmao, alphao):
147     """
148         Method which performs selection. The best mu parents are selected from the
149         pool of lambd offspring (, selection), or from the combined pool of
150         previous mu parents and their lambd offspring (+ selection).
151
152     Parameters
153     -----
154     xp : (mu, n) array
155         Matrix of mu previous parent solutions, each of dimension n.
156     fxp : (mu, ) array
157         Vector of function value evaluations for the mu previous parent
158         solutions.
159     sigmap : (mu, n) array
160         Matrix of standard deviations for the previous parent solutions.
161     alphap : (mu, n, n) array
162         Matrix of rotation angles for the previous parent solutions.
163     xo : (lambd, n) array
164         Matrix of n-dimensional offspring solutions.
165     fxo : (lambd, ) array
166         Vector of function value evaluations for the lambd offspring
167         solutions.
168     sigmao : (lambd, n) array
169         Matrix of standard deviations for the offpsring solutions.
170     alphao : (lambd, n, n) array
171         Matrix of rotation angles for the offspring solutions.
172
173     Returns
174     -----
175     xp : (mu, n) array
176         Matrix of mu parent solutions, each of dimension n.
177     fxp : (mu, ) array
178         Vector of function value evaluations for the mu parent solutions.
179     sigmap : (mu, n) array
180         Matrix of standard deviations for the parent solutions.
181     alphap : (mu, n, n) array
182         Matrix of rotation angles for the parent solutions.
183     """
184
185     if self.sel == ',':
186         # select the mu lowest solutions from the offspring only
187         idx = np.argsort(fxo)[:self.mu]
188         xp = xo[idx]
189         fxp = fxo[idx]
190         sigmap = sigmao[idx]
191         alphap = alphao[idx]
192         return xp, fxp, sigmap, alphap
193     else:
194         # append the parents and offspring solutions to form a combined pool
195         # for selection
196         xc = np.append(xo, xp, axis=0)

```

```

191     fxc = np.append(fxo, fxp, axis=0)
192     sigmac = np.append(sigmo, sigmap, axis=0)
193     alphac = np.append(alphao, alphap, axis=0)
194     # select the mu lowest solutions from the combined pool
195     idx = np.argsort(fxc)[:, :self.mu]
196     xc = xc[idx]
197     fxp = fxc[idx]
198     sigmap = sigmac[idx]
199     alphap = alphac[idx]
200     return xc, fxp, sigmap, alphap
201
202 def mutation(self, x, sigma, alpha):
203     """
204         Method which performs mutation of the control variables and strategy
205         parameters.
206
207     Parameters
208     -----
209     x : (mu, n) array
210         Matrix of mu parent solutions, each of dimension n.
211     sigma : (mu, n) array
212         Matrix of parent standard deviations.
213     alpha : (mu, n, n)
214         Matrix of parent rotation angles.
215
216     Returns
217     -----
218     xm : (mu, n) array
219         Matrix of mu mutated solutions, each of dimension n.
220     sigmam : (mu, n) array
221         Matrix of mutated standard deviations.
222     alpham : (mu, n, n)
223         Matrix of mutated rotation angles.
224
225     # number of parents
226     num_p = x.shape[0]
227     # realisation of a N(0, 1) Gaussian, used in the mutation of all sigma
228     # values
229     chi_0 = np.random.randn()
230     # construct a (mu, n, n) matrix of random numbers drawn from N(0, 1)
231     # Gaussian distributions
232     randn = np.random.randn(num_p, self.n, self.n)
233     # construct a (mu, n, n) matrix of chi's from the realisation tmp
234     # the upper triangular of tmp is taken
235     # the lower triangular is constructed as the transpose of the upper
236     # triangular of tmp (excluding the leading diagonal)
237     # matrix is thus symmetric in axis 1 and 2 (chi_ij = chi_ji)
238     chi = np.triu(randn) + np.transpose(np.triu(randn, 1), (0, 2, 1))
239     # the (mu, n) matrix of chi_i's is taken as the diagonal elements of the
240     # (mu, n, n) matrix chi over axis 1 and 2
241     chi_i = np.diagonal(chi, axis1=1, axis2=2)
242     # mutate the standard deviations
243     sigmam = sigma * np.exp(self.tau * chi_i + self.tau_d * chi_0)
244     # mutate the rotation angles
245     alpham = alpha + self.beta * chi

```

```

241 # compute the off-diagonal elements of the (mu, n, n) mutated covariance
242 # matrix
243 # the first bracketed term is a (mu, n, n) matrix of variance differences
244 # of which the ij-th term is the difference ( $\sigma_{\text{m}, i}^2 - \sigma_{\text{m}, j}^2$ )
245 # it is constructed by adding a new axis to each of the sigma matrices
246 # such that they are (mu, n, 1) and transposing one matrix over axis 1 and
247 # 2
248 # the second bracketed term is a (mu, n, n) matrix of two times the
249 # tangent of the mutated rotation angles
250 # the two bracketed terms are multiplied element-wise to give the (mu, n,
251 # n) mutated covariance matrix
252 covm = 0.5 * (sigmam[:, :, None] ** 2 - np.transpose(sigmam[:, :, None],
253 (0, 2, 1)) ** 2) * np.tan(2 * alpham)
254 # get the shape (mu, n, n) of the mutated covariance matrix
255 s0, s1, s2 = covm.shape
256 # set the diagonal elements over axis 1 and 2 of the (mu, n, n) mutated
257 # covariance matrix to the mutated variances
258 # a small term epsilon_c is added to ensure the mutated covariance matrix
259 # is PSD
260 covm.reshape(s0, -1)[:, ::s2 + 1] = sigmam ** 2 + self.epsilon_c
261 # draw a (mu, n) matrix of perturbations dx from Gaussian distributions
262 # dx[i] (n, ) is drawn from a  $N(0, \text{covm}[i])$  distribution
263 dx = np.array([np.random.multivariate_normal(mean=np.zeros(self.n),
264 cov=covmi) for covmi in covm])
265 # mutate the solution x by applying the perturbation dx
266 xm = x + dx
267 return xm, sigmam, alpham
268
269
270 def discrete_recombination(self, xp, no):
271 """
272     Method which discretely recombines parent solutions to form a given number
273     of offspring.
274
275     Parameters
276     -----
277     xp : (mu, n) array
278         Matrix of mu parent solutions, each of dimension n.
279     no : int
280         Number of offspring to generate.
281
282     Returns
283     -----
284     xo : (no, n) array
285         Matrix of no offspring solutions, each of dimension n.
286     """
287 num_p = xp.shape[0]
288 # randomly select parents 1 and 2 from the pool of potential parents
289 x1 = xp[np.random.randint(low=0, high=num_p, size=no)]
290 x2 = xp[np.random.randint(low=0, high=num_p, size=no)]
291 # construct a binary mask for parent 1
292 x1_mask = np.random.randint(low=0, high=2, size=np.array(x1.shape))
293 # parent 2's mask is complementary to parent 1's
294 x2_mask = np.logical_not(x1_mask)
295 # recombine using masks to form offspring
296 xo = x1_mask * x1 + x2_mask * x2

```

```

287         return xo
288
289     def intermediate_recombination(self, sigmap, alphap, no):
290         """
291             Method which intermediately recombines parent strategy parameters to form
292             a given number of offspring strategy parameters.
293
294         Parameters
295         -----
296         sigmap : (mu, n) array
297             Matrix of parent standard deviations.
298         alphap : (mu, n, n) array
299             Matrix of parent rotation angles.
300         no : int
301             Number of offspring to generate.
302
303         Returns
304         -----
305         sigmap : (mu, n) array
306             Matrix of offspring standard deviations.
307         alphap : (mu, n, n) array
308             Matrix of offspring rotation angles.
309         """
310
311         num_p = sigmap.shape[0]
312         # randomly select parents 1 and 2 from the pool of potential parents
313         idx1 = np.random.randint(low=0, high=num_p, size=no)
314         idx2 = np.random.randint(low=0, high=num_p, size=no)
315
316         sigma1 = sigmap[idx1]
317         sigma2 = sigmap[idx2]
318         # offspring standard deviation is formed as the arithmetic mean of the
319         # parent standard deviations
320         sigmao = 0.5 * (sigma1 + sigma2)
321         # force standard deviations to exceed a minimum value
322         sigmao = np.where(sigmao > self.epsilon_s, sigmao, self.epsilon_s)
323         # prevent standard deviations from exceeding a maximum value
324         sigmao = np.where(sigmao < 100, sigmao, 100)
325
326         alpha1 = alphap[idx1]
327         alpha2 = alphap[idx2]
328         # offspring rotation angles are formed as the arithmetic mean of the
329         # parent rotation angles
330         alphao = 0.5 * (alpha1 + alpha2)
331         # circularly map rotation angles which lie outside the range [-pi, pi] by
332         # adding/subtracting a factor of 2 * pi
333         alphao = np.where(alphao < np.pi, alphao, alphao - 2 * np.pi)
334         alphao = np.where(alphao > - np.pi, alphao, alphao + 2 * np.pi)
335         return sigmao, alphao
336
337     def update_archive(self, archive_x, archive_fx, x_j, f_j):
338         """
339             Method which updates the archive of best solutions.
340
341         Parameters
342         -----

```

```

338     archive_x : (k, n) array
339         Matrix of k solutions, each of n dimensions, in the current archive.
340     archive_fx : (k, ) array
341         Vector of the objective function evaluated at the k solutions in the
342         current archive.
343     x_j : (n, ) array
344         A new n-dimensional candidate solution for archiving.
345     f_j : float
346         Objective function evaluated at the candidate solution x_j.

347     Returns
348     -----
349     archive_x : (l, n) array
350         Matrix of k solutions, each of n dimensions, in the updated archive.
351     archive_fx : (l, ) array
352         Vector of the objective function evaluated at the l solutions in the
353         updated archive.
354     """
355     # worst archived solution
356     g_ind = np.argmax(archive_fx)
357     f_g = archive_fx[g_ind]

358     # Euclidean distance between new and archived solutions
359     distance = norm(archive_x - x_j, axis=1)
360     # archived solution x_e which most closely resembles new solution x_j
361     # (closest in Euclidean distance)
362     e_ind = np.argmin(distance)
363     d_ej = distance[e_ind]
364     f_e = archive_fx[e_ind]

365     # archive is not full
366     if len(archive_x) < self.archive_size:
367         # archive new solution x_j if it is sufficiently dissimilar to all the
368         # solutions archived
369         if d_ej > self.d_min:
370             archive_x.append(x_j)
371             archive_fx.append(f_j)

372     # archive is full
373     else:
374         # archive x_j if it is sufficiently dissimilar to all the solutions
375         # archived and better than the worst
376         if d_ej > self.d_min and f_j < f_g:
377             # new solution x_j replaces the worst archived solution x_g
378             archive_x.pop(g_ind)
379             archive_fx.pop(g_ind)
380             archive_x.append(x_j)
381             archive_fx.append(f_j)
382         # archive x_j if it is not sufficiently dissimilar to all the
383         # solutions archived and better than the worst
384         if d_ej < self.d_min and f_j < f_g:
385             # new solution x_j replaces the closest archived solution x_e
386             archive_x.pop(e_ind)
            archive_fx.pop(e_ind)
            archive_x.append(x_j)

```

```

387             archive_fx.append(f_j)
388             # archive  $x_j$  if it is not the best solution so far and sufficiently
389             # similar to but better than  $x_e$ 
390             if d_ej < self.d_sim and f_j < f_e:
391                 archive_x.pop(e_ind)
392                 archive_fx.pop(e_ind)
393                 archive_x.append(x_j)
394                 archive_fx.append(f_j)
395             return archive_x, archive_fx
396
397     def termination(self, fxp):
398         """
399             Method which tests the termination condition.
400
401             Parameters
402             -----
403             fxp : (mu, ) array
404                 Vector of objective function evaluations for mu parent solutions.
405
406             Returns
407             -----
408             bool
409                 True if termination condition is met. False otherwise.
410
411             # test if absolute difference between worst and best solutions lies below
412             # the threshold epsilon
413             return np.abs(max(fxp) - min(fxp)) < self.epsilon
414
415     def run(self):
416         """
417             Method which runs the ES algorithm.
418
419             Returns
420             -----
421             x_min : (n, ) array
422                 Minimised solution.
423             fx_min : float
424                 Objective function evaluated at the minimised solution.
425             archive_x : (k, n) array
426                 Matrix of k final archived solutions. Each control variable is of
427                 dimension n.
428             archive_fx : (k, ) array
429                 Vector of objective function values evaluated at the k final archived
430                 solutions.
431
432             # generate initial solutions and control parameters
433             xp, fxp, sigmap, alphap, iters = self.initialise()
434             # form the initial archives from the initial solutions
435             archive_x = [xp[0]]
436             archive_fx = [fxp[0]]
437             for i in range(xp.shape[0]):
438                 archive_x, archive_fx = self.update_archive(archive_x, archive_fx,
439                     xp[i], fxp[i])
440
441             while iters < self.max_iter:

```

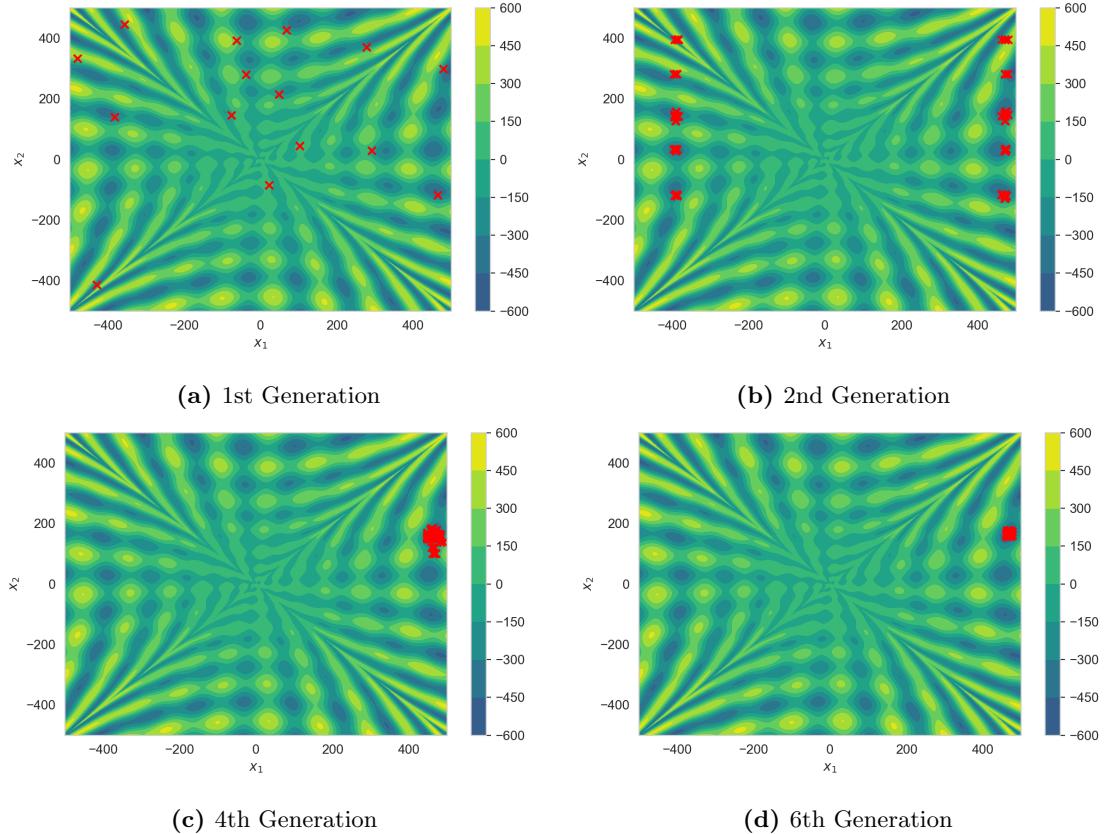
```

437     if self.termination(fxp):
438         break
439     # initialise array of accepted offspring
440     xo, sigmao, alphao = [], [], []
441     # generate lambda valid offspring
442     while len(xo) < self.lambd:
443         no = self.lambd - len(xo)
444         xm, sigmam, alpham = self.mutation(xp, sigmap, alphap)
445         x_n = self.discrete_recombination(xm, no)
446         sigma_n, alpha_n = self.intermediate_recombination(sigmam, alpham,
447                                                 no)
448         # keep offspring solutions which adhere to inequality constraints
449         idx = tuple([np.max(np.abs(x_n), axis=1) < self.bound])
450         x_n, sigma_n, alpha_n = x_n[idx], sigma_n[idx], alpha_n[idx]
451         # append valid offspring solutions to array of accepted offspring
452         xo = np.append(xo, x_n).reshape(len(xo) + len(x_n), self.n)
453         sigmao = np.append(sigmao, sigma_n).reshape(len(sigmao) +
454                                                 len(sigma_n), self.n)
455         alphao = np.append(alphao, alpha_n).reshape(len(alphao) +
456                                                 len(alpha_n), self.n, self.n)
457
458         fxo = self.f(xo)
459         iters += self.lambd
460         xp, fpx, sigmap, alphap = self.selection(xo, fxo, sigmao, alphao, xp,
461                                                fpx, sigmap, alphap)
462
463         # update archive with the parent solutions
464         for i in range(xp.shape[0]):
465             archive_x, archive_fx = self.update_archive(archive_x, archive_fx,
466                                              xp[i], fpx[i])
467
468         # best minimised solution
469         idx = np.argmin(archive_fx)
470         x_min = archive_x[idx]
471         fx_min = archive_fx[idx]
472
473         return x_min, fx_min, archive_x, archive_fx
474
475 def main():
476     # number of dimensions and bound on feasible region
477     n, bound = 5, 500
478     es = ES(f, n, bound)
479     x_min, fx_min, archive_x, archive_fx = es.run()

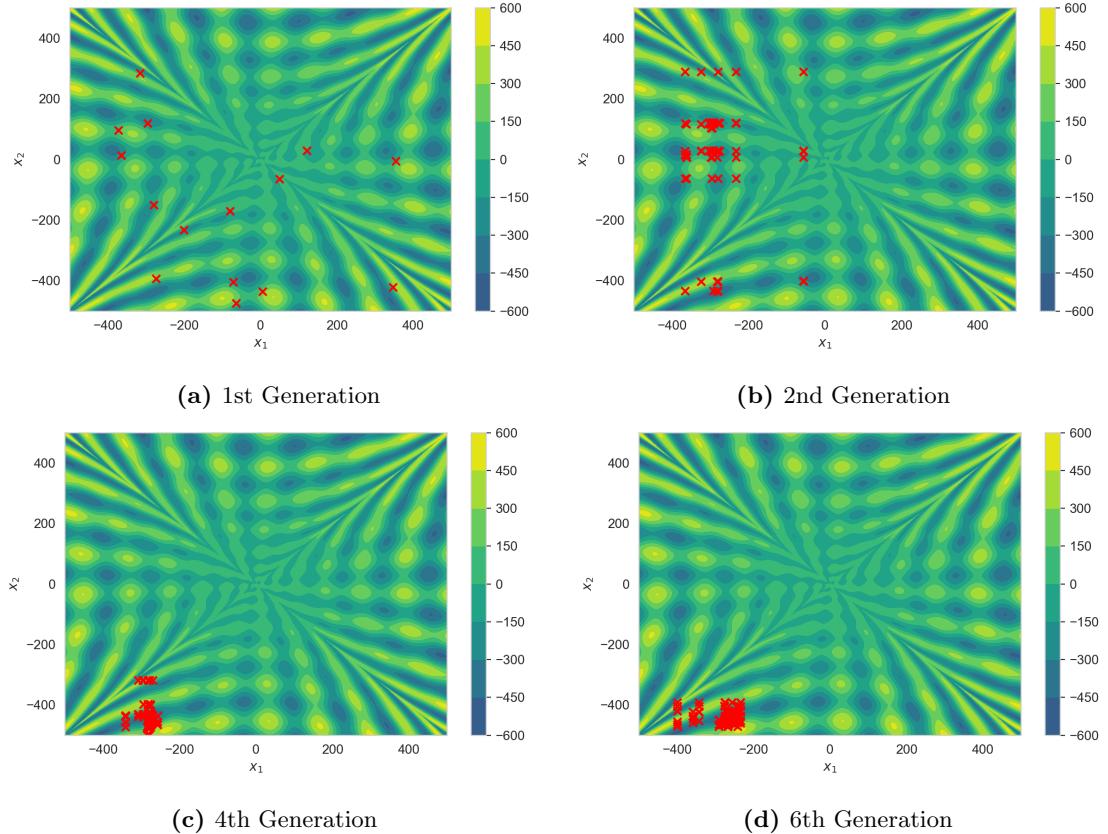
```

---

## B Evolution Strategies Population Evolution



**Figure 16:** Evolution of the population distribution for the minimisation of the 2D-RF using the ES algorithm with the default parameter configuration using seed 1. The algorithm converges to a local minimum at (467, 165).



**Figure 17:** Evolution of the population distribution for the minimisation of the 2D-RF using the ES algorithm with the default parameter configuration using seed 2. The algorithm converges to a local minimum at (-281, -440).